

insurance

June 10, 2020

```
[1]: import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
data = pd.read_csv('datasets_13720_18513_insurance.csv')
data.head()
```

```
[1]:   age    sex    bmi  children  smoker    region    charges
0   19  female  27.900         0     yes  southwest  16884.92400
1   18   male  33.770         1     no   southeast   1725.55230
2   28   male  33.000         3     no   southeast   4449.46200
3   33   male  22.705         0     no  northwest  21984.47061
4   32   male  28.880         0     no  northwest   3866.85520
```

```
[2]: data.isna().sum()
```


```
[2]: age      0
sex        0
bmi        0
children   0
smoker     0
region     0
charges    0
dtype: int64
```

```
[3]: f= plt.figure(figsize=(12,5))

ax=f.add_subplot(121)
sns.distplot(data[(data.smoker == 'yes')]['charges'],color='c',ax=ax)
ax.set_title('Distribution of charges for smokers')

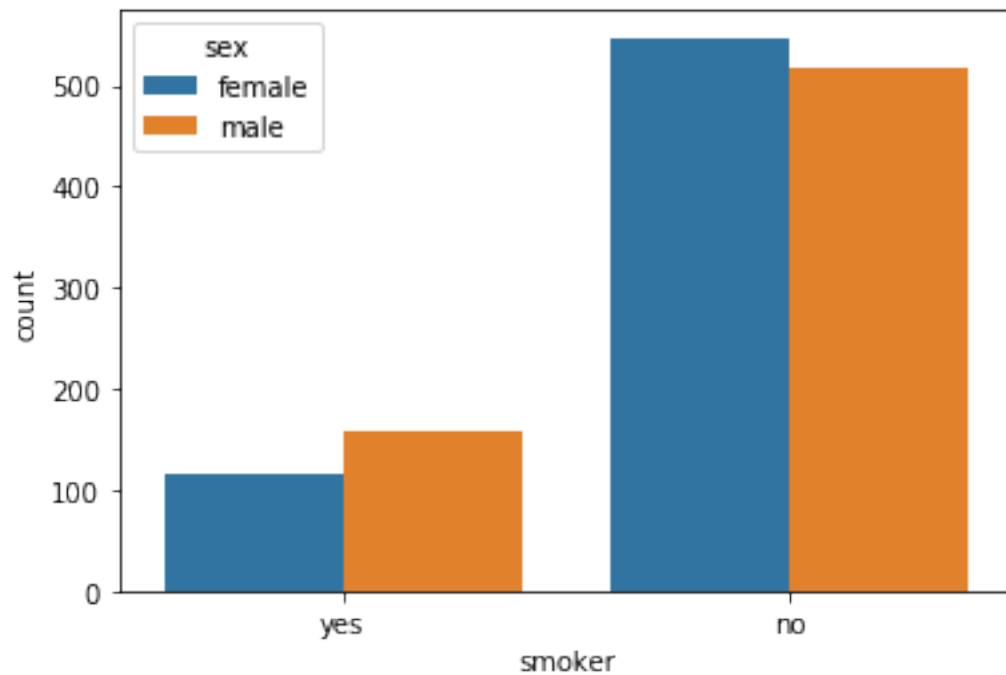
ax=f.add_subplot(122)
sns.distplot(data[(data.smoker == 'no')]['charges'],color='b',ax=ax)
ax.set_title('Distribution of charges for non-smokers')
```

```
[3]: Text(0.5, 1.0, 'Distribution of charges for non-smokers')
```




```
[4]: sns.countplot(x = 'smoker', hue = 'sex', data = data)
```

```
[4]: <matplotlib.axes._subplots.AxesSubplot at 0x1a169fdeb8>
```



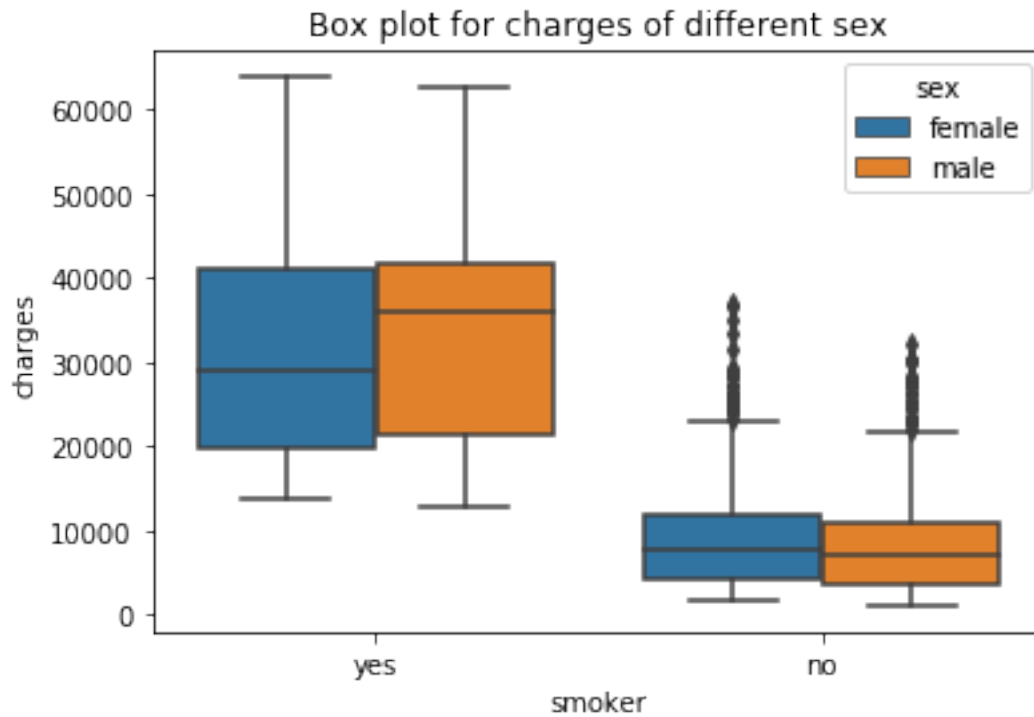
```
[5]: sns.catplot(x = 'sex', y = 'charges', hue = 'smoker', data = data, kind = 'violin')
```

```
[5]: <seaborn.axisgrid.FacetGrid at 0x1a16b9bf28>
```




```
[6]: plt.title("Box plot for charges of different sex")
sns.boxplot(x = 'smoker', y = 'charges', hue = 'sex', data= data)
```

```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16c90198>
```



```
[7]: plt.figure(figsize=(12,5))
plt.title("Distribution of age")
sns.distplot(data['age'])
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16d69a20>
```




```
[8]: data[data['age']==18]['smoker'].value_counts()
```

```
[8]: no      57  
     yes     12  
     Name: smoker, dtype: int64
```


```
[9]: sns.countplot(x = 'smoker', hue = 'sex', data = data[data['age'] == 18])
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16e523c8>
```



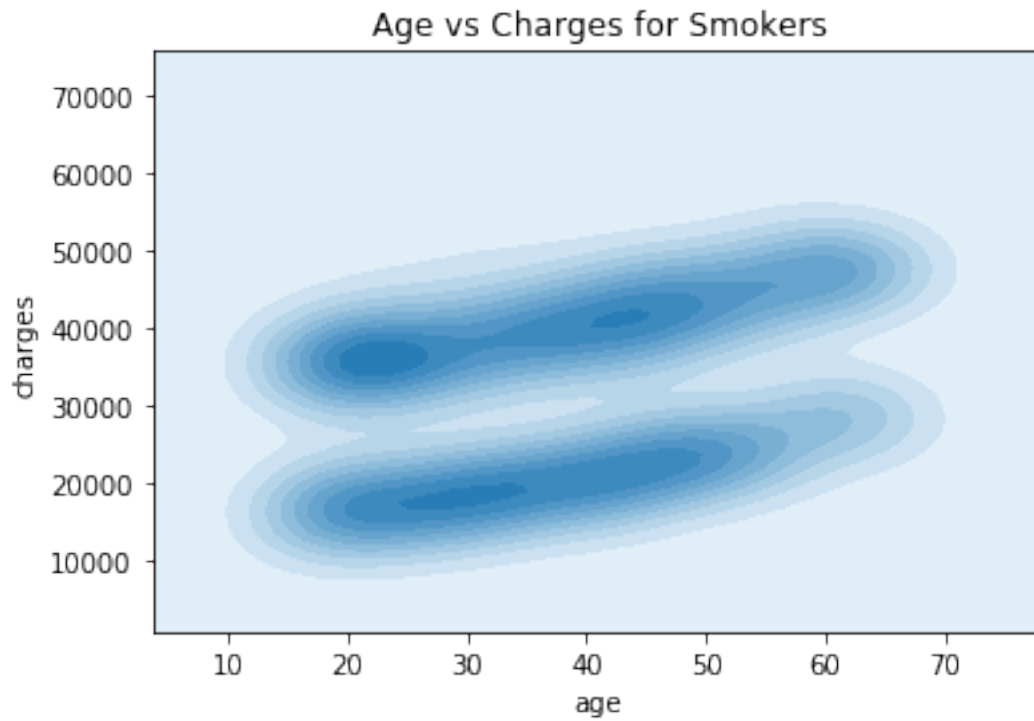
```
[10]: # Does smoking affect the cost of treatment at age 18?  
      plt.figure(figsize = (5, 8))  
      sns.boxplot(x = 'smoker', y = 'charges', data = data[data['age'] == 18])
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16e739e8>
```



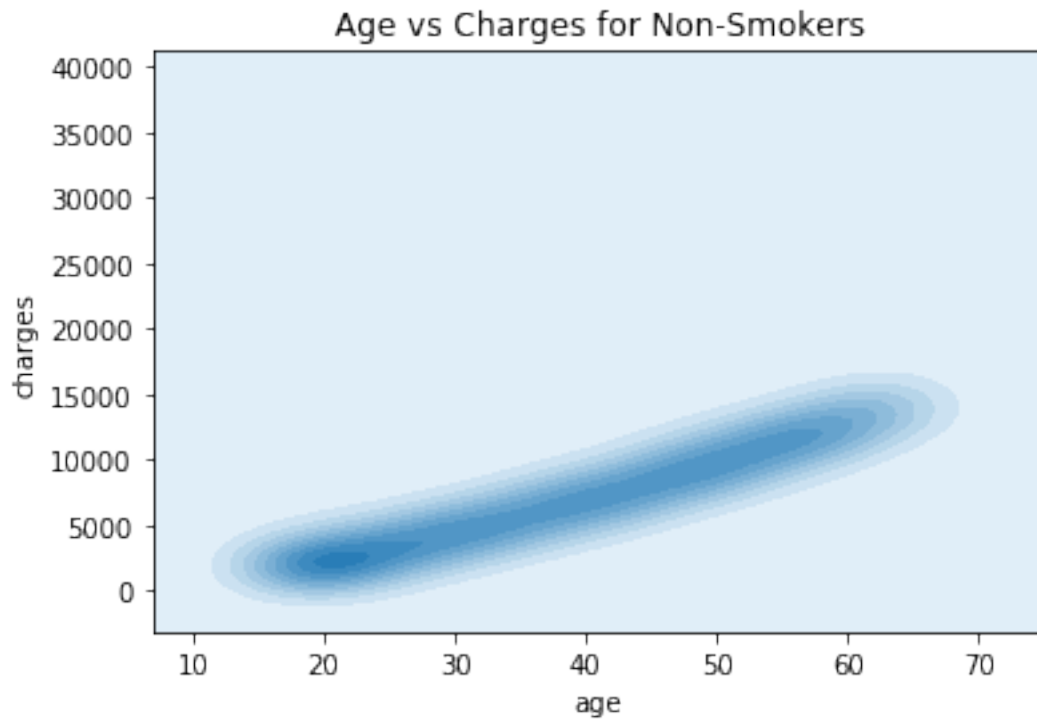
```
[11]: sns.kdeplot(data[data['smoker'] == 'yes']['age'],data[data['smoker'] == 'no']['age'], shade = True)  
plt.title('Age vs Charges for Smokers')
```

```
[11]: Text(0.5, 1.0, 'Age vs Charges for Smokers')
```



```
[12]: sns.kdeplot(data[data['smoker'] == 'no']['age'],data[data['smoker'] == 'no']['charges'], shade = True)  
plt.title('Age vs Charges for Non-Smokers')
```


```
[12]: Text(0.5, 1.0, 'Age vs Charges for Non-Smokers')
```




```
[14]: plt.figure(figsize=(10,10))
      sns.lmplot(x = 'age', y = 'charges', hue = 'smoker', data = data, markers = ['o', 'x'])
```

```
[14]: <seaborn.axisgrid.FacetGrid at 0x1a176114e0>
```

<Figure size 720x720 with 0 Axes>

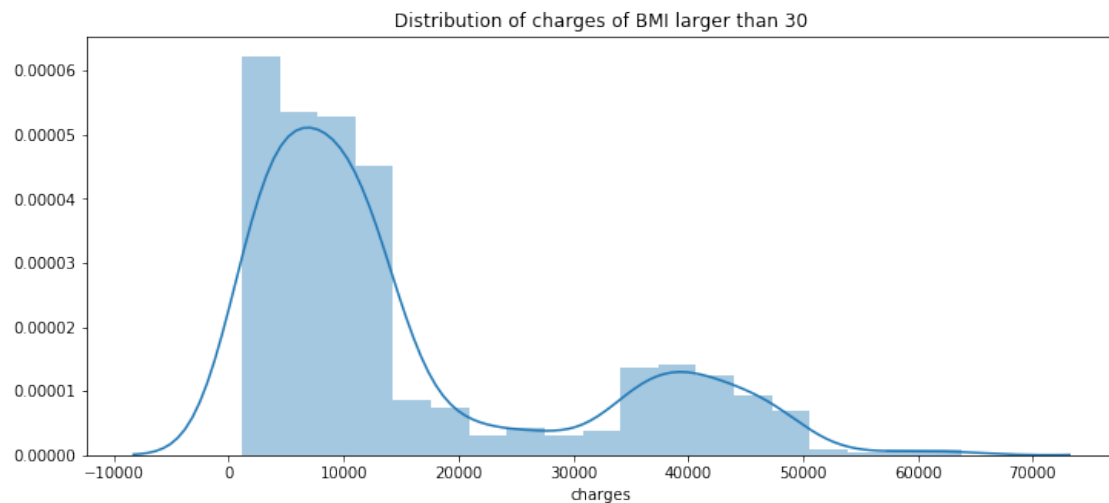


```
[17]: ### bmi distribution  
plt.figure(figsize=(12,5))  
plt.title("Distribution of bmi")  
ax = sns.distplot(data["bmi"])
```



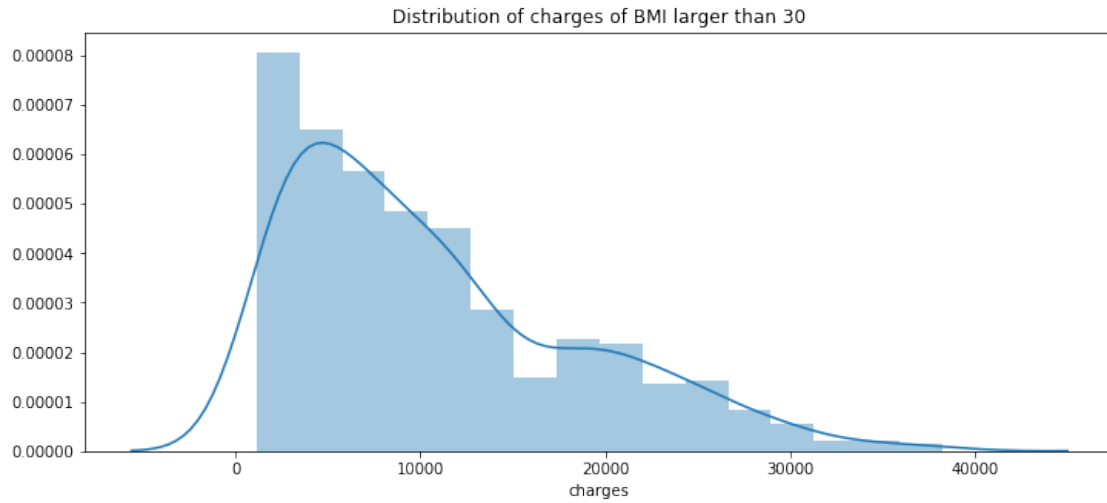
```
[19]: ### bmi score of 30 is a cutoff point
plt.figure(figsize = (12, 5))
plt.title('Distribution of charges of BMI larger than 30')
sns.distplot(data[data['bmi'] >= 30]['charges'])
```

[19]: <matplotlib.axes._subplots.AxesSubplot at 0x1a19252e80>




```
[20]: plt.figure(figsize = (12, 5))
plt.title('Distribution of charges of BMI smaller than 30')
sns.distplot(data[data['bmi'] < 30]['charges'])
```

[20]: <matplotlib.axes._subplots.AxesSubplot at 0x1a18c39c50>




```
[23]: sns.kdeplot(data['bmi'], data['charges'], shade = True)
```

```
[23]: <matplotlib.axes._subplots.AxesSubplot at 0x1a18d0aef0>
```




```
[24]: sns.jointplot(x="bmi", y="charges", data = data, kind="kde", color="r")
```

```
[24]: <seaborn.axisgrid.JointGrid at 0x1a190243c8>
```




```
[27]: ### bmi score vs charges smoker and non-smoker  
plt.figure(figsize = (9, 6))  
sns.scatterplot(x = 'bmi', y = 'charges', hue = 'smoker', data = data)
```

```
[27]: <matplotlib.axes._subplots.AxesSubplot at 0x1a197bfd30>
```




```
[29]: sns.countplot(data['children'])
```

```
[29]: <matplotlib.axes._subplots.AxesSubplot at 0x1a19b314e0>
```



```
[30]: sns.boxplot(x = 'children', y = 'charges', data = data)
```

```
[30]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1a3b2908>
```



```
[31]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
le.fit(data.sex)
data['sex'] = le.transform(data['sex'])

le.fit(data.smoker)
data.smoker = le.transform(data.smoker)

le.fit(data.region)
data.region = le.transform(data.region)

plt.figure(figsize = (10, 9))
sns.heatmap(data.corr(),annot = True, fmt = '.3f')
```

```
[31]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1a1d7a20>
```



```
[32]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.ensemble import RandomForestRegressor
```

```
[42]: X = data.drop(['charges'], axis = 1)
y = data['charges']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 0)
```

```
[45]: lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
print('r2 train:', lr.score(X_train, y_train))
print('r2 test:', lr.score(X_test, y_test))
```

```
print(mean_squared_error(y_test, y_pred))
```

```
r2 train: 0.7337162219022217  
r2 test: 0.7962732059725786  
32073628.56010921
```

```
[46]: n = list(range(2, 6))  
r2_train = []  
r2_test = []  
mse = []  
for i in range(2, 6):  
    quad = PolynomialFeatures(degree = i)  
    x_quad = quad.fit_transform(X)  
    X_train, X_test, y_train, y_test = train_test_split(x_quad, y, random_state=  
    ↳ 0)  
  
    plr = LinearRegression()  
    plr.fit(X_train, y_train)  
  
    y_pred = plr.predict(X_test)  
    r2_train.append(plr.score(X_train, y_train))  
    r2_test.append(plr.score(X_test, y_test))  
    mse.append(mean_squared_error(y_test, y_pred))  
result = pd.DataFrame(np.column_stack([n, r2_train, r2_test, mse]),  
                        columns=['Degree', 'train Rsquared', 'test_  
↳Rsquared', 'Mse'])  
result
```

```
[46]:
```

	Degree	train Rsquared	test Rsquared	Mse
0	2.0	0.831481	0.884628	1.816348e+07
1	3.0	0.841715	0.879056	1.904082e+07
2	4.0	0.856473	0.857891	2.237286e+07
3	5.0	0.884269	0.781857	3.434326e+07

```
[65]: X = data.drop(['charges'], axis = 1)  
y = data['charges']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 0)
```

```
[73]: rf = RandomForestRegressor(oob_score = True, random_state = 1, criterion= 'mse')  
rf.fit(X_train, y_train)  
  
y_pred = rf.predict(X_test)  
  
print(r2_score(y_test, y_pred))  
print(mean_squared_error(y_test, y_pred))
```

```
0.8707069513939738  
20355188.114508193
```



```
[61]: #print(rf.oob_score_)
```

```

└─
└─-----
AttributeError                                Traceback (most recent call
└─last)

<ipython-input-61-07119920495b> in <module>
----> 1 print(rf.oob_score_)

AttributeError: 'RandomForestRegressor' object has no attribute
└─'oob_score_'
```

```
[51]: importances = rf.feature_importances_
std = np.std([tree.feature_importances_ for tree in rf.estimators_],
             axis=0)
indices = np.argsort(importances)[::-1]

# Print the feature ranking
print("Feature ranking:")

for f in range(X.shape[1]):
    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))

# Plot the feature importances of the forest
plt.figure()
plt.title("Feature importances")
plt.bar(range(X.shape[1]), importances[indices],
        color="r", yerr=std[indices], align="center")

plt.xticks(range(X_train.shape[1]), X.columns[indices], rotation=90)

plt.xlim([-1, X.shape[1]])
plt.show()
```

Feature ranking:

1. feature 4 (0.603241)
2. feature 2 (0.213779)
3. feature 0 (0.138449)
4. feature 3 (0.021509)
5. feature 5 (0.016149)
6. feature 1 (0.006872)

